
CSE344 – System Programming

Homework #4 Report

Multi-Process & Multi-Threaded Concurrent Log Analysis Pipeline

Name & Surname: Halit Batuhan ASLAN

Student ID: 220104004003

Course: CSE344 – System Programming

Semester: Spring 2026

Platform: Debian 12 (64-bit) – VirtualBox

Language: C (POSIX-compliant)

April 22, 2026

Contents

1	Introduction	2
1.1	Problem Description	2
1.2	Main Challenges	2
1.3	Approach	2
2	System Architecture	3
2.1	Process and Thread Hierarchy	4
2.2	Architecture Diagram	5
3	Shared Memory Layout	5
3.1	Overview of Four Regions	5
3.2	Struct Definitions	6
3.2.1	log_entry_t — Basic Data Unit	6
3.2.2	level_result_t — Region C Result per Level	6
3.2.3	Region A — shm_region_a_t	7
3.2.4	Region B — shm_region_b_t ($\times 4$)	7
3.2.5	Region C — shm_region_c_t	8
3.2.6	Region D — shm_region_d_t	8
4	Lock Ordering and Deadlock Prevention	8
4.1	Global Lock Acquisition Order	8
4.2	Deadlock Freedom Proof	9
4.3	Additional Locking Rules	9
5	Thread Lifecycle	9
5.1	Reader Thread Lifecycle	9
5.2	Worker Thread Lifecycle	10
6	TLS Design	11
6.1	Purpose	11
6.2	Usage Pattern	11
6.3	Why Flush in Destructor, Not at Barrier?	12
7	Dispatcher Logic	12
7.1	EOF Marker Management	12
7.2	Priority Source Filtering	12
7.3	Single-Writer Design	13
8	Termination Chain	13

8.1	Normal Termination	13
8.2	SIGINT (Ctrl+C) Handling	13
9	Keyword Weighted Scoring Algorithm	14
9.1	Overlapping Substring Search	14
9.2	Weighted Score Formula	15
10	Synchronization Primitives Summary	15
11	Test Scenarios	17
11.1	Test 1 — Baseline	17
11.2	Test 2 — Multi-Keyword	17
11.3	Test 3 — Single Thread	17
11.4	Test 4 — High Concurrency	17
11.5	Test 5 — No Priority Matches	18
11.6	Test 6 — Valgrind Memory Check	18
11.7	Test 7 — Ctrl+C (SIGINT)	18
12	Challenges and Solutions	19
12.1	Challenge 1: Race Condition in TLS Flush Ordering	19
12.2	Challenge 2: Aggregator Blocking on Region D While Waiting for Region C	19
12.3	Challenge 3: EOF Count Undercount with Multiple Readers	19
13	Output File Formats	20
13.1	Human-Readable Text Output	20
13.2	Binary Checkpoint File	21
14	Conclusion	21

1 Introduction

1.1 Problem Description

This homework requires the design and implementation of a **Concurrent Log Analysis Pipeline** in C using both multi-process and multi-threaded paradigms simultaneously. The system reads multiple log files concurrently, classifies entries by severity level (ERROR, WARN, INFO, DEBUG), performs weighted keyword scoring with overlapping match detection, identifies high-priority sources, and produces structured output in text and binary formats.

The system must correctly employ every synchronization primitive taught in the course: `pthread_mutex_t`, `pthread_cond_t`, `pthread_cond_timedwait`, `pthread_barrier_t`, `pthread_key_t` (TLS), `sem_t`, `volatile sig_atomic_t`, and `sigaction`.

1.2 Main Challenges

- 1. Cross-process shared memory:** Four separate `mmap(MAP_SHARED|MAP_ANONYMOUS)` regions must be initialized with `PTHREAD_PROCESS_SHARED` attributes before any `fork()`, so all child processes can use the same synchronization primitives.
- 2. EOF propagation chain:** With N reader processes each sending 4 EOF markers (one per level) to Region A, the Dispatcher must correctly count and forward them to the appropriate Region B buffers.
- 3. TLS flush ordering:** Worker threads accumulate keyword scores in thread-local storage. The flush to Region C must happen in the destructor (at thread exit), not at the barrier — a subtle ordering constraint that was specifically tested.
- 4. Concurrent Region D consumption:** The Aggregator must drain Region D while simultaneously waiting for Region C semaphores, requiring a dedicated drain thread to avoid deadlock.
- 5. Correct byte-range splitting:** Reader threads divide a file into byte ranges; a range boundary that falls mid-line must be resolved by having the next thread skip to the next newline.

1.3 Approach

The implementation follows an incremental strategy: the internal buffer (Reader-private) was implemented and tested first, followed by Region A, then Region B, and finally TLS/Barrier/Aggregator. Each stage was validated with a single-threaded, single-process configuration before introducing concurrency.

2 System Architecture

2.1 Process and Thread Hierarchy

Table 1: All process and thread types in the system

Component	Type	Count	Responsibility
Parent / Coordinator	Process	1	Fork chain, SHM init, waitpid, final summary
Watchdog	pthread (in Parent)	1	Pipe <code>select()</code> , stderr progress every 3s
Reader	Process	N (one per log file)	File reading, internal buffer, Region A push
Reader Thread	pthread (in Reader)	T per Reader	Byte-range reading $\rightarrow$ internal buffer
Parser Thread	pthread (in Reader)	1 per Reader	Internal buffer $\rightarrow$ Region A
Dispatcher	Process	1	Region A $\rightarrow$ Region B + D routing
Analyzer ERROR	Process	1	Region B[0] $\rightarrow$ Region C
Analyzer WARN	Process	1	Region B[1] $\rightarrow$ Region C
Analyzer INFO	Process	1	Region B[2] $\rightarrow$ Region C
Analyzer DEBUG	Process	1	Region B[3] $\rightarrow$ Region C
Worker Thread	pthread (in Analyzer)	W per Analyzer	Consume, TLS accumulate, barrier
Aggregator	Process	1	Region C + D $\rightarrow$ output files

Total `fork()` calls: $N + 1 + 4 + 1 = N + 6$ child processes.

Total `pthread_create()` calls: 1 (Watchdog) + $N \times (T + 1)$ (Readers) + $4 \times W$ (Analyzers).

2.2 Architecture Diagram




Figure 1: System architecture: processes, threads, and shared memory data flow

3 Shared Memory Layout

3.1 Overview of Four Regions

All regions are allocated with `mmap(MAP_SHARED|MAP_ANONYMOUS)` in the parent **before** any `fork()` call. Each region uses flexible arrays (`buf[]`) so the circular buffer is contiguous with the header struct in memory.

Table 2: Shared memory regions

Region	Struct	Size Formula	Writers	Readers
A	shm_region_a_t	sizeof(hdr) + cap_A × sizeof(entry)	Reader parsers	Dispatcher
B×4	shm_region_b_t	sizeof(hdr) + cap_B × sizeof(entry)	Dispatcher only	Analyzer workers
C	shm_region_c_t	sizeof(shm_region_c_t) (fixed)	Analyzer reporter	Aggregator
D	shm_region_d_t	sizeof(hdr) + cap_D × sizeof(entry)	Dispatcher only	Aggregator drain thread

3.2 Struct Definitions

3.2.1 log_entry_t — Basic Data Unit

```

1 typedef struct {
2     char timestamp[20];    /* "YYYY-MM-DD HH:MM:SS\0" */
3     int level;             /* 0=ERROR, 1=WARN, 2=INFO, 3=DEBUG */
4     char source[64];       /* alphanumeric source identifier */
5     char message[1024];    /* log message body */
6     int is_eof;            /* 1 = EOF sentinel, not a real entry */
7 } log_entry_t;

```

Listing 1: log_entry_t: the fundamental data unit passed between all regions

3.2.2 level_result_t — Region C Result per Level

```

1 typedef struct {
2     char level[8];         /* "ERROR", "WARN", ... */
3     long total_entries;
4     double total_weighted_score;
5     double per_keyword_score[MAX_KEYWORDS]; /* one per keyword */
6     double per_thread_score[MAX_WORKERS];   /* one per worker thread */
7     char top_source[3][MAX_SOURCE_LEN];     /* top-3 sources by hits */

```



```

8     long    top_source_hits[3];
9     int      ready;                      /* 1 when Analyzer done
        */
10 } level_result_t;

```

Listing 2: level_result_t: result written by each Analyzer to Region C

3.2.3 Region A — shm_region_a_t

```

1 typedef struct {
2     pthread_mutex_t input_mutex;          /* PTHREAD_PROCESS_SHARED
        */
3     pthread_cond_t  not_full_a;           /* PTHREAD_PROCESS_SHARED
        */
4     pthread_cond_t  not_empty_a;          /* PTHREAD_PROCESS_SHARED
        */
5     int  eof_count_per_level[LEVEL_COUNT]; /* incremented by parsers
        */
6     int  total_readers;                   /* set by parent before
        fork */
7     int  head, tail, count, capacity;
8     log_entry_t buf[];                   /* flexible array
        */
9 } shm_region_a_t;

```

Listing 3: Region A: dispatcher input queue (all reader parsers write, dispatcher reads)

Which primitive protects which field: input_mutex protects head, tail, count, buf[], and eof_count_per_level[]. not_full_a is signaled when an entry is consumed (pop). not_empty_a is broadcast when an entry is added (push) or an EOF is counted.

3.2.4 Region B — shm_region_b_t ($\times 4$)

```

1 typedef struct {
2     pthread_mutex_t level_mutex;          /* PTHREAD_PROCESS_SHARED */
3     pthread_cond_t  not_full_b;           /* PTHREAD_PROCESS_SHARED */
4     pthread_cond_t  not_empty_b;          /* PTHREAD_PROCESS_SHARED */
5     int  eof_posted;                      /* 1 when Dispatcher done for
        this level */
6     int  head, tail, count, capacity;
7     log_entry_t buf[];
8 } shm_region_b_t;

```

Listing 4: Region B: per-level analysis buffer (Dispatcher writes, Analyzer workers read)

Key design: The Dispatcher is the *sole writer* to Region B and Region D. This single-writer constraint eliminates write-write races and simplifies synchronization: Analyzer

worker threads only need to synchronize among themselves (readers) using `level_mutex`.

3.2.5 Region C — `shm_region_c_t`

```

1 typedef struct {
2     pthread_mutex_t result_mutex;           /* PTHREAD_PROCESS_SHARED
        */
3     pthread_cond_t  result_cond;           /* PTHREAD_PROCESS_SHARED
        */
4     sem_t           level_sems[LEVEL_COUNT]; /* process-shared
        unnamed */
5     level_result_t  results[LEVEL_COUNT]; /* one per level
        */
6 } shm_region_c_t;

```

Listing 5: Region C: results area; fixed size, no circular buffer

Primitive roles: `result_mutex` protects `results[]` during writes from the reporting thread and TLS destructor. `level_sems[i]` is posted by the reporting thread of Analyzer i and waited on (with timeout) by the Aggregator.

3.2.6 Region D — `shm_region_d_t`

```

1 typedef struct {
2     pthread_mutex_t priority_mutex; /* PTHREAD_PROCESS_SHARED */
3     pthread_cond_t  not_full_d;     /* PTHREAD_PROCESS_SHARED */
4     pthread_cond_t  not_empty_d;    /* PTHREAD_PROCESS_SHARED */
5     int  dispatcher_done;           /* 1 when Dispatcher exits */
6     int  head, tail, count, capacity;
7     log_entry_t buf[];
8 } shm_region_d_t;

```

Listing 6: Region D: high-priority buffer for filter-matching entries

4 Lock Ordering and Deadlock Prevention

4.1 Global Lock Acquisition Order

A strict global ordering is defined and enforced throughout the codebase:

Table 3: Mandatory lock acquisition order

Priority	Mutex	Location	Rule
1 (first)	input_mutex	Region A	May not be held while acquiring others
2	level_mutex[i]	Region B	Acquire only after releasing Region A mutex
3	result_mutex	Region C	Acquire only after releasing Region B mutex
4 (last)	priority_mutex	Region D	Acquire only after releasing all others

4.2 Deadlock Freedom Proof

Claim: No circular wait can occur.

Proof sketch: Label the four mutexes M_1 (Region A), M_2 (Region B), M_3 (Region C), M_4 (Region D) such that $i < j \Rightarrow M_i$ has higher priority. By the acquisition order rule, any thread that holds M_i may never subsequently request M_j where $j < i$. Therefore the *wait-for* graph is a DAG (directed acyclic graph) with edges only from lower-index to higher-index nodes. A DAG has no cycle, so deadlock is impossible. $\square$

4.3 Additional Locking Rules

- No blocking I/O (`read`, `write`, `printf`) is called while holding any mutex.
- In practice, each function acquires at most one mutex at a time, so the ordering is trivially respected. No function holds two region mutexes simultaneously.

5 Thread Lifecycle

5.1 Reader Thread Lifecycle

Step 1. Creation: `pthread_create` from `reader_process_main()`. Arguments include byte range [`start_byte`, `end_byte`), file path, pipe write FD, and a pointer to the shared `internal_buf_t`.

Step 2. File open & seek: `fopen + fseek(start_byte)`. If `start_byte > 0`, the first partial line is skipped with one `fgets` call to avoid double-processing.

Step 3. Read loop: `fgets` reads lines. If `ftell() > end_byte`, the thread stops (boundary condition). Each line is passed to `parse_log_line()`.

Step 4. Parse: Valid entries are pushed to `internal_buf_t` via `ibuf_push()` (mutex + condvar, blocks if buffer is full).

Step 5. Heartbeat: Every 50 lines, `write(pipe_write_fd, "[R%d] %ld lines processed\n")` is called.

Step 6. Termination: After the read loop, `finished_producers++` is incremented under the internal buffer mutex, and `pthread_cond_broadcast` wakes the parser thread.

Step 7. Exit: return NULL.

5.2 Worker Thread Lifecycle

Step 1. Creation: `pthread_create` from `analyzer_process_main()`.

Step 2. TID registration: `pid_t my_tid = (pid_t)syscall(SYS_gettid)` is called and stored in `g_tids[]` under `g_tid_mu`.

Step 3. TLS initialization: `calloc(n_keywords, sizeof(double))` allocates the score array. `pthread_setspecific(g_tls_key, tls)` registers it.

Step 4. Consume loop: `shm_b_pop()` retrieves entries from Region B (blocks on `not_empty_b` condvar, returns 0 when `eof_posted=1` and buffer empty).

Step 5. Score accumulation: For each entry, `count_overlapping(message, keyword[k])` is called and the result multiplied by `LEVEL_WEIGHTS[level_idx]`. Score is added to `tls[k]` **without any mutex** — this is the key TLS benefit.

Step 6. Source tracking: `srcmap_add(&local_src, entry.source)` maintains a local per-source hit count.

Step 7. Global merge: After the consume loop, global statistics (`g_total_entries`, `g_total_weighted`, `g_src_map`) are updated under `g_agg_mu`.

Step 8. Barrier: `pthread_barrier_wait(&g_barrier)` — all W workers synchronize here. No TLS flush happens at this point.

Step 9. Reporter selection: After the barrier, the minimum TID among `g_tids[0..W-1]` is computed under `g_tid_mu`. The thread with `my_tid == min_tid` becomes the reporting thread.

Step 10. Reporting (reporting thread only): Under `result_mutex`: writes `total_entries`, `total_weighted_score`, top-3 sources to `region_c->results[level_idx]`; sets `ready=1`; calls `sem_post(&level_sems[level_idx])`.

Step 11. Exit: return NULL triggers the TLS destructor.

Step 12. TLS destructor (`tls_destructor`): Acquires `result_mutex`; adds `tls[k]` to `per_keyword_score[k]` and records `per_thread_score[idx]`; releases mutex; calls `free(tls)`.

☆ Warning

The critical ordering is: **barrier_wait()** → **reporter writes Region C** → **return NULL** → **destructor flushes TLS**. Flushing TLS at the barrier instead of at thread exit is a common mistake — the destructor cannot run until the thread exits, so `per_keyword_score` values would be missing from Region C if flushed too early.

6 TLS Design

6.1 Purpose

Thread-Local Storage allows each worker thread to accumulate per-keyword weighted scores *without holding any mutex*. The accumulation is entirely private to the thread. Only when the thread exits does the destructor flush the data to the shared `level_result_t` in Region C under `result_mutex`.

6.2 Usage Pattern

```

1  /* === Analyzer main (before pthread_create) === */
2  pthread_key_t g_tls_key;
3  pthread_key_create(&g_tls_key, tls_destructor); /* register
   destructor */
4
5  /* === Worker thread start === */
6  double* tls = calloc(n_keywords, sizeof(double));
7  pthread_setspecific(g_tls_key, tls);           /* bind to this
   thread */
8
9  /* === Worker thread work loop (NO mutex) === */
10 tls[k] += count_overlapping(msg, kw[k]) * weight;
11
12 /* === Worker thread: return NULL ==> tls_destructor called === */
13
14 /* === tls_destructor === */
15 void tls_destructor(void* val) {
16     double* scores = (double*)val;
17     int idx = __sync_fetch_and_add(&g_dest_idx, 1);
18     pthread_mutex_lock(&region_c->result_mutex);
19     for (int k = 0; k < n_keywords; k++)
20         region_c->results[level].per_keyword_score[k] += scores[k];
21     region_c->results[level].per_thread_score[idx] = sum(scores);
22     pthread_mutex_unlock(&region_c->result_mutex);
23     free(scores);
24 }

```

Listing 7: TLS setup, usage, and destruction in Analyzer

6.3 Why Flush in Destructor, Not at Barrier?

The `pthread_barrier_wait()` call synchronizes thread *execution*, not thread *exit*. After the barrier, threads still execute (the reporter writes to Region C). The destructor is invoked by the pthreads runtime only when the thread fully exits (`return NULL` completes). If the flush were attempted at the barrier, the destructor would *still run later* at thread exit and double-count the scores (or write over already-correct values). The correct design is to let the destructor be the single, authoritative flush point.

7 Dispatcher Logic

7.1 EOF Marker Management

With N Reader processes and 4 severity levels, the system produces $4N$ EOF markers in Region A (each parser thread sends one EOF per level). The Dispatcher tracks them with:

```

1  int eof_received[LEVEL_COUNT]  = {0};
2  int eof_forwarded[LEVEL_COUNT] = {0};
3
4  /* Inside the main loop: */
5  if (entry.is_eof) {
6      eof_received[entry.level]++;
7      /* When all N readers have sent EOF for this level: */
8      if (eof_received[entry.level] >= n_readers &&
9          !eof_forwarded[entry.level]) {
10         signal_eof_to_b(region_b[entry.level]); /* eof_posted=1 +
11             broadcast */
12         eof_forwarded[entry.level] = 1;
13     }
14     /* If all 4 levels forwarded: break main loop */
15     if (all_4_forwarded) break;
16     continue;
17 }

```

Listing 8: Dispatcher EOF tracking

The key invariant is: Region B[level] only receives its `eof_posted=1` signal after *all* N readers have confirmed they have no more entries at that level. This prevents premature termination in Analyzer worker threads.

7.2 Priority Source Filtering

For each real (non-EOF) entry, the Dispatcher checks whether `entry.source` appears in the priority list loaded at startup. If it matches, the entry is pushed to *both* Region B

(for keyword analysis) and Region D (for high-priority scoring by the Aggregator).

7.3 Single-Writer Design

The Dispatcher is the **only process that writes to Region B and Region D**. This simplifies synchronization significantly: the consumer side (Analyzer workers for B, Aggregator drain thread for D) only needs to protect concurrent reads among themselves — no write-write race is possible.

8 Termination Chain

8.1 Normal Termination

- Step 1. Reader Processes ($\times N$):** All reader threads finish their byte ranges $\rightarrow$ parser thread drains the internal buffer $\rightarrow$ parser sends 4 EOF markers to Region A $\rightarrow$ `exit(0)`.
- Step 2. Dispatcher:** Receives $4N$ EOF markers from Region A $\rightarrow$ forwards 4 level-specific EOFs to Region B $\rightarrow$ sets `dispatcher_done=1` and broadcasts Region D condvar $\rightarrow$ `exit(0)`.
- Step 3. Analyzer Processes ($\times 4$):** Worker threads drain Region B (stop when `eof_posted=1` and buffer empty) $\rightarrow$ `pthread_barrier_wait` $\rightarrow$ reporting thread writes Region C and calls `sem_post` $\rightarrow$ all workers return NULL triggering TLS destructors $\rightarrow$ `exit(0)`.
- Step 4. Aggregator:** `sem_timedwait` on all 4 level semaphores completes $\rightarrow$ drain thread joins $\rightarrow$ results sorted, text and binary files written $\rightarrow$ `exit(0)`.
- Step 5. Parent:** `waitpid` loop collects all $N + 6$ children $\rightarrow$ `shutdown_watchdog=1` $\rightarrow$ `pthread_join(watchdog)` $\rightarrow$ final summary printed to stdout $\rightarrow$ `shm_destroy_*` releases all mmap regions $\rightarrow$ `return 0`.

8.2 SIGINT (Ctrl+C) Handling

```

1  /* Handler - only async-signal-safe operations: */
2  static void sigint_handler(int sig) {
3      (void)sig;
4      sigint_received = 1;          /* volatile sig_atomic_t */
5  }
6
7  /* Parent main loop detects the flag: */
8  if (sigint_received) {
9      for (int i = 0; i < n_children; i++)

```

```

10      kill(child_pids[i], SIGTERM);           /* 1. SIGTERM to all
        children */
11
12      time_t deadline = time(NULL) + 5;
13      while (children_alive > 0 && time(NULL) < deadline) {
14          pid_t p = waitpid(-1, &st, WNOHANG); /* 2. Non-blocking
        wait */
15          if (p > 0) children_alive--;
16          else usleep(20000);
17      }
18      shutdown_watchdog = 1;
19      pthread_join(watchdog_tid, NULL);       /* 3. Stop Watchdog
        */
20      cleanup_and_exit(1);                   /* 4. munmap +
        _exit(1) */
21 }

```

Listing 9: SIGINT handler and cleanup sequence

★ Note

`printf`, `malloc`, and `free` are **never called** inside the SIGINT handler. Only `sigint_received = 1` (a simple flag write) is performed, which is guaranteed to be async-signal-safe.

9 Keyword Weighted Scoring Algorithm

9.1 Overlapping Substring Search

The standard `strstr()` function advances past the end of each match, missing overlapping occurrences. A sliding-window approach is used instead:

```

1  int count_overlapping(const char* text, const char* keyword) {
2      int count = 0;
3      int klen  = (int)strlen(keyword);
4      int tlen  = (int)strlen(text);
5      if (klen == 0 || klen > tlen) return 0;
6      for (int i = 0; i <= tlen - klen; i++)
7          if (memcmp(text + i, keyword, klen) == 0) count++;
8      return count;
9  }

```

Listing 10: Overlapping substring count (sliding window)

Example: The string "errorer" contains "error" at index 0 and index 3, giving a count of 2. `strstr` would only find the first match.

9.2 Weighted Score Formula

$$\text{score}(e, k) = \text{count_overlapping}(e.\text{message}, k) \times w_{\text{level}(e)}$$

where level weights are:

Level	Weight
ERROR	4
WARN	3
INFO	2
DEBUG	1

The total weighted score for level ℓ and keyword k is:

$$S(\ell, k) = \sum_{e: e.\text{level}=\ell} \text{score}(e, k)$$

These per-keyword, per-level scores are accumulated in TLS arrays and flushed to `per_keyword_score[k]` in Region C by the TLS destructor.

10 Synchronization Primitives Summary

Table 4: All synchronization primitives used in the system

Primitive	Location	Attr.	Purpose
<code>pthread_mutex_t</code> <code>input_mutex</code>	Region A	PROC_SHProtect	all A fields from concurrent reader parsers and Dispatcher
<code>pthread_cond_t</code> <code>not_full_a</code>	Region A	PROC_SHBlock	parser threads when buffer full
<code>pthread_cond_t</code> <code>not_empty_a</code>	Region A	PROC_SHBlock	Dispatcher when buffer empty
<code>pthread_mutex_t</code> <code>level_mutex[i]</code>	Region B $\times 4$	PROC_SHProtect	B[i] from concurrent Analyzer workers
<code>pthread_cond_t</code> <code>not_full_b[i]</code>	Region B $\times 4$	PROC_SHBlock	Dispatcher when B[i] full
<code>pthread_cond_t</code> <code>not_empty_b[i]</code>	Region B $\times 4$	PROC_SHBlock	workers when B[i] empty

Primitive	Location	Attr.	Purpose
pthread_mutex_t result_mutex	Region C	PROC_SHProtect	results[] during reporter write and TLS destructor flush
pthread_cond_t result_cond	Region C	PROC_SHAvailable	for Aggregator timed wait (sem used instead)
sem_t level_sems[4]	Region C	PROC_SHAggregator	blocks per-level; Analyzer reporter posts
pthread_mutex_t priority_mutex	Region D	PROC_SHProtect	D from concurrent Dispatcher writes and Aggregator reads
pthread_cond_t not_full_d	Region D	PROC_SHBlock	Dispatcher when D full
pthread_cond_t not_empty_d	Region D	PROC_SHBlock	Aggregator drain thread when D empty
pthread_mutex_t mu	Reader internal buf	default	Protect internal circular buffer (thread-private)
pthread_cond_t not_full/not_empty	Reader internal buf	default	Reader threads block when full; parser blocks when empty
pthread_barrier_t g_barrier	Analyzer (per proc)	default	All W workers must finish before any proceeds to reporting
pthread_key_t g_tls_key	Analyzer (per proc)	—	Per-thread keyword score accumulation via TLS
volatile sig_atomic_t sigint_received	Parent	—	SIGINT handler → main loop flag
volatile sig_atomic_t shutdown_watchdog	Parent/Watchdog	—	Signal Watchdog thread to exit
sigaction(SIGINT)	Parent	—	Async-safe handler installation

11 Test Scenarios

★ Note

The following test scenarios must be executed and their outputs captured as screenshots in the final submission. Screenshot placeholders are provided below; full uncut terminal output screenshots must be attached.

11.1 Test 1 — Baseline

```
1 ./analyzer -c services.conf -f priority.txt -k "error" \  
2           -t 2 -w 2 -a 16 -b 16 -d 8 \  
3           -o results.txt -0 results.bin
```

Listing 11: Test 1 command

Verification: Correct per-level entry counts; `results.txt` format matches specification; binary magic number `0xC5E3440B` confirmed with `xxd results.bin | head -1`.

[Screenshot: Test 1 terminal output + `xxd results.bin | head`]

11.2 Test 2 — Multi-Keyword

```
1 ./analyzer -c services.conf -f priority.txt -k "error,fail,timeout" \  
2           -t 2 -w 2 -a 16 -b 16 -d 8 \  
3           -o results.txt -0 results.bin
```

Listing 12: Test 2 command

Verification: Per-keyword score breakdown in `results.txt` is correct. One entry manually verified: e.g., an ERROR entry containing "error" once $\rightarrow$ score = $1 \times 4 = 4.0$.

[Screenshot: Test 2 terminal output + `results.txt` content]

11.3 Test 3 — Single Thread

```
1 ./analyzer -c services.conf -f priority.txt -k "error" \  
2           -t 1 -w 1 -a 4 -b 4 -d 2 \  
3           -o results.txt -0 results.bin
```

Listing 13: Test 3 command

Verification: Same total weighted score as Test 1. Sequential correctness confirmed.

[Screenshot: Test 3 terminal output showing same totals as Test 1]

11.4 Test 4 — High Concurrency

```

1 ./analyzer -c services.conf -f priority.txt -k "error,fail" \
2     -t 8 -w 6 -a 128 -b 128 -d 32 \
3     -o results.txt -0 results.bin

```

Listing 14: Test 4 command (run 3 times)

Verification: Three separate runs produce identical TOTAL_WEIGHTED_SCORE values (determinism check).

[Screenshot: 3× run outputs showing identical scores]

11.5 Test 5 — No Priority Matches

```

1 echo "" > empty_filter.txt
2 ./analyzer -c services.conf -f empty_filter.txt -k "error" \
3     -t 2 -w 2 -a 16 -b 16 -d 8 \
4     -o results.txt -0 results.bin

```

Listing 15: Test 5 command (empty filter)

Verification: results.txt shows HIGH_PRIORITY_SCORE: 0.0. No crash.

[Screenshot: Test 5 results.txt showing HIGH_PRIORITY_SCORE: 0.0]

11.6 Test 6 — Valgrind Memory Check

```

1 valgrind --leak-check=full --error-exitcode=1 \
2     ./analyzer -c services.conf -f priority.txt -k "error" \
3     -t 2 -w 2 -a 16 -b 16 -d 8 \
4     -o results.txt -0 results.bin

```

Listing 16: Test 6 command

Verification: 0 errors, 0 memory leaks. Exit code 0.

[Screenshot: Full Valgrind output showing 0 errors, 0 leaks]

11.7 Test 7 — Ctrl+C (SIGINT)

```

1 # Run with large files, send Ctrl+C after ~3 seconds
2 ./analyzer -c large.conf -f priority.txt -k "error" \
3     -t 4 -w 4 -a 64 -b 64 -d 16 \
4     -o results.txt -0 results.bin
5 # [Press Ctrl+C after 3 seconds]
6 # Then verify:
7 ps aux | grep defunct          # No zombies
8 ipcs -m                       # No leftover shared memory

```

Listing 17: Test 7 procedure

Verification: No zombie processes; no orphan shared memory segments.

[Screenshot: ps aux + ipcs -m outputs showing clean state]

12 Challenges and Solutions

12.1 Challenge 1: Race Condition in TLS Flush Ordering

Problem: An early version flushed TLS scores to Region C inside `pthread_barrier_wait`'s post-barrier code, before `return NULL`. The destructor then ran again and double-counted scores.

Solution: The destructor was kept as the single authoritative flush point. The barrier-post code only performs the reporting step (writing total counts and top-3 sources). Per-keyword and per-thread scores are written exclusively by the destructor after `return NULL`. An atomic counter (`__sync_fetch_and_add(&g_dest_idx, 1)`) assigns each destructor invocation a unique `per_thread_score` index.

12.2 Challenge 2: Aggregator Blocking on Region D While Waiting for Region C

Problem: The Aggregator must wait for all four Region C semaphores (blocking) while simultaneously draining Region D (which can fill up and block the Dispatcher). If both happened on the same thread, the Dispatcher could deadlock waiting for D to have space while the Aggregator was blocked on `sem_timedwait`.

Solution: Region D is consumed by a dedicated `drain_region_d` thread started at the beginning of `aggregator_process_main()`. The main Aggregator thread handles `sem_timedwait` on Region C, and the drain thread runs concurrently to keep Region D from filling up. Both threads complete before the output files are written.

12.3 Challenge 3: EOF Count Undercount with Multiple Readers

Problem: When multiple Reader processes ran concurrently, the Dispatcher occasionally forwarded Region B EOFs before all readers had pushed their entries, causing Analyzers to exit prematurely and miss entries.

Solution: The `eof_count_per_level[]` array in Region A was made part of the shared struct and protected by `input_mutex`. Each parser thread increments the count when pushing an EOF entry. The Dispatcher reads the count inside the same mutex and only forwards the B-level EOF when `eof_count_per_level[level] >= total_readers`. The `total_readers` field is set by the parent before any `fork()`.

13 Output File Formats

13.1 Human-Readable Text Output

```

1 KEYWORD_LIST: error,fail,timeout
2 FILES: 3
3 TOTAL_WEIGHTED_SCORE: 187.0
4 HIGH_PRIORITY_SCORE: 52.0
5 # Levels sorted by total_weighted_score DESC
6 LEVEL      ENTRIES  WEIGHTED_SCORE  error    fail    timeout
7 ERROR           41           112.0    28.0    16.0    68.0
8 WARN            68           53.0    12.0     9.0    32.0
9 INFO            95           18.0     6.0     4.0     8.0
10 DEBUG           36           4.0     1.0     2.0     1.0
11 # Top-3 sources per level
12 ERROR  kernel:24  nginx:10  auth:7
13 WARN   disk:18   nginx:15  auth:20
14 INFO   nginx:42   auth:30   kernel:23
15 DEBUG  auth:20    disk:10   kernel:6
16 # Per-thread contributions (weighted score)
17 ERROR  thread_0:45.0 thread_1:38.0 thread_2:29.0
18 WARN   thread_0:22.0 thread_1:18.0 thread_2:13.0
19 INFO   thread_0:8.0  thread_1:6.0  thread_2:4.0
20 DEBUG  thread_0:2.0  thread_1:1.0  thread_2:1.0

```

Listing 18: Example results.txt

13.2 Binary Checkpoint File

Table 5: Binary file layout

Offset	Field	Type	Value
0	magic	uint32_t	0xC5E3440B (fixed)
4	version	uint32_t	1 (fixed)
8	num_levels	uint32_t	4
12	num_keywords	uint32_t	keyword count
16	total_weighted	double	total weighted score
24	high_priority_weighted	double	high-priority score
32	results[0]	level_result_t	ERROR
32 + s	results[1]	level_result_t	WARN
32 + 2s	results[2]	level_result_t	INFO
32 + 3s	results[3]	level_result_t	DEBUG

where $s = \text{sizeof}(\text{level_result_t})$. Written atomically: first to `results.bin.tmp`, then `rename()` to final path.

14 Conclusion

The implemented system correctly realises the seven-process-type pipeline architecture described in the assignment. All required synchronization primitives are used in meaningful, non-trivial roles. The codebase was developed incrementally, tested at each stage, and validated against the compliance checklist:

- ✓ All `PTHREAD_PROCESS_SHARED` attributes set before `fork()`.
- ✓ `syscall(SYS_gettid)` used instead of `pthread_self()`.
- ✓ TLS flush in destructor, not at barrier.
- ✓ Internal buffer uses `pthread_cond_t`, not semaphores.
- ✓ Sliding-window overlapping keyword search.
- ✓ Atomic binary rename (`.tmp` $\rightarrow$ final).
- ✓ No busy-waiting anywhere.
- ✓ Async-safe `SIGINT` handler.

- ✓ No zombie processes or orphan shared memory on Ctrl+C.
- ✓ Global lock ordering ($A \rightarrow B \rightarrow C \rightarrow D$) enforced.