

CSE 344 – System Programming

Spring 2026

Homework #4 Report

*Multi-Process & Multi-Threaded
Concurrent Log Analysis Pipeline*

Student Name: Halit Batuhan ASLAN

Student ID: 220104004003

Course: CSE 344 – System Programming

Homework: #4

Due Date: 02 April 2026

Platform: Debian 11 (64-bit) on VirtualBox

Compiler: gcc -Wall -Wextra -pthread -O2

Submitted in fulfillment of the requirements of CSE 344.

May 1, 2026

Contents

1	Introduction	2
1.1	Problem Description	2
1.2	Main Challenges	2
1.3	Approach	2
2	System Architecture	2
2.1	Process and Thread Inventory	2
2.2	Architecture Diagram	3
2.3	Communication Channels	4
3	Shared Memory Layout	4
3.1	Region Allocation	4
3.2	Region A – Dispatcher Input Queue	4
3.3	Region B – Per-Level Buffers ($\times 4$)	4
3.4	Region C – Result Area	5
3.5	Region D – Priority Shadow Buffer	5
4	Lock Ordering and Deadlock Freedom	6
4.1	Mandatory Acquisition Order	6
4.2	Lock Order Graph	6
4.3	Proof Sketch – No Circular Wait	6
5	Thread Lifecycle	7
5.1	Reader Thread	7
5.2	Worker Thread	7
5.3	Reporting Thread Selection	8
6	Thread-Local Storage Design	8
6.1	Why TLS at All?	8
6.2	Implementation	8
6.3	Why Flush in the Destructor and Not at the Barrier?	9
6.4	Synchronizing the Reporting Thread with Destructors	10
7	Dispatcher Logic	10
7.1	Routing	10
7.2	Multi-Reader EOF Handling	10
7.3	Termination Detection	11
8	Termination Chain	11
8.1	Normal Termination	11
8.2	SIGINT (Ctrl+C) Termination	11
9	Test Scenarios	11
9.1	Test 1 – Baseline	12
9.2	Test 2 – Multi-keyword Weighted Scoring	14
9.3	Test 3 – Single Thread	16
9.4	Test 4 – High Concurrency & Determinism	18
9.5	Test 5 – Empty Filter	18
9.6	Test 6 – Valgrind (Memory Leaks)	20
9.7	Test 7 – Ctrl+C Cleanup	20

10 Challenges and Solutions	21
10.1 Challenge 1: TLS Destructor Double-Flush	21
10.2 Challenge 2: Region A Sliding-Window Termination	21
10.3 Challenge 3: Atomic Binary Write	21
11 Build and Run Instructions	22
11.1 Compilation	22
11.2 Sample Run	22
12 Conclusion	22

1 Introduction

1.1 Problem Description

The objective of this assignment is to design and implement a **Concurrent Log Analysis Pipeline** that exercises every major synchronization primitive taught throughout the semester. The system reads multiple log files concurrently, classifies log entries by severity level (ERROR, WARN, INFO, DEBUG), counts occurrences of multiple user-specified keywords with level-weighted scoring, identifies high-priority sources via a separate filtering stage, and produces both a human-readable report and a binary checkpoint file.

The pipeline is organized as a sequence of cooperating processes: *Reader* processes read log files; a *Dispatcher* routes entries by level; four *Analyzer* processes perform per-level keyword counting; an *Aggregator* writes the final outputs; a *Watchdog* thread inside the parent prints periodic progress snapshots. Communication happens exclusively through four shared memory regions (A, B, C, D) protected by process-shared synchronization primitives.

1.2 Main Challenges

This assignment is significantly more challenging than a textbook producer-consumer problem. The principal difficulties are:

- **Two paradigms in one program.** Multi-process and multi-thread concurrency must coexist correctly. Process-shared primitives must use the `PTHREAD_PROCESS_SHARED` attribute, whereas the in-process buffer inside each Reader uses default (private) primitives.
- **Multi-keyword weighted scoring with overlapping matches.** "errorerror" contains "error" two times (indices 0 and 3). A naive use of `strstr` would miss this; a proper sliding-window matcher is required.
- **Thread-Local Storage with destructor flush.** Each worker thread must accumulate its scores in TLS and rely on a `pthread_key_create` destructor to flush them at thread exit. Flushing at the barrier (instead of in the destructor) is a common mistake and was specifically warned against.
- **Timed waits with deterministic termination.** `pthread_cond_timedwait` and `sem_timedwait` must be used so the pipeline can detect EOF deterministically and never hang on a missing producer.
- **Atomic binary write.** The binary checkpoint file must be written to a `.tmp` file and atomically renamed via `rename(2)`, so that a partially-written file can never appear at the final path.
- **Lock ordering.** A strict global order ($A \rightarrow B \rightarrow C \rightarrow D$) must be defined and respected to guarantee deadlock freedom.

1.3 Approach

The implementation follows a layered design. `shm.c/shm.h` provide the four shared regions and bounded-buffer primitives; each component (*reader*, *dispatcher*, *analyzer*, *aggregator*, *watchdog*) is then a thin layer on top. `main.c` orchestrates the lifecycle: it allocates every shared region *before* forking, forks all seven components, spawns the Watchdog thread, waits on every child, prints the system summary, then destroys everything. SIGINT triggers a safe broadcast to all children with a five-second deadline before the parent forcibly cleans up.

2 System Architecture

2.1 Process and Thread Inventory

The running system consists of **seven distinct process/thread types**:

1. **Parent / Coordinator** (1 process) – forks all children, runs the `waitpid` loop, prints the final summary.
2. **Reader Process** (one per log file) – internally spawns T reader threads plus one parser thread.
3. **Dispatcher Process** (1 process) – routes entries from Region A into Region B[level] and Region D.
4. **Analyzer Process** (4 processes, one per severity level) – internally spawns W worker threads; the lowest-TID worker becomes the reporting thread.
5. **Aggregator Process** (1 process) – writes the final `.txt` and `.bin` outputs; spawns one helper thread to drain Region D.
6. **Watchdog Thread** (1 thread inside the parent) – `select()`-based pipe multiplexer; prints progress to `stderr` every three seconds.
7. **Reader Thread / Worker Thread / Parser Thread** (many threads inside Reader and Analyzer processes).

2.2 Architecture Diagram




Figure 1: System architecture. Solid arrows show the data path through the shared memory regions; dashed arrows show heartbeat pipes used by the Watchdog. The four Region B sub-buffers act as parallel lanes, one per severity level.

2.3 Communication Channels

Channel	From	To	Purpose
Region A	Reader (parser thread)	Dispatcher	all parsed entries flow through this single ring buffer
Region B[i]	Dispatcher	Analyzer[i]	per-level entries (one buffer per level)
Region C	Analyzer reporting thread	Aggregator	one <code>level_result_t</code> per level
Region D	Dispatcher	Aggregator drainer	copies of entries with priority sources
Pipe[i]	Reader[i] (heartbeat)	Watchdog (parent)	ASCII heartbeat lines for progress display

Table 1: Inter-component communication channels.

3 Shared Memory Layout

3.1 Region Allocation

All four regions are allocated by the parent *before any fork()* using anonymous shared mmap:

```
void* alloc_shared(size_t sz)
{
    void* p = mmap(NULL, sz, PROT_READ | PROT_WRITE,
                   MAP_SHARED | MAP_ANONYMOUS, -1, 0);
    if (p == MAP_FAILED) { perror("mmap"); exit(1); }
    memset(p, 0, sz);
    return p;
}
```

Children inherit the mappings across `fork()`, so every process sees identical bytes. Synchronization primitives stored *inside* these regions are initialized with `PTHREAD_PROCESS_SHARED`, otherwise their behavior across process boundaries is undefined.

3.2 Region A – Dispatcher Input Queue

Single circular buffer of `log_entry_t` (capacity from `-a`). Written by every Reader’s parser thread; read by the Dispatcher.

```
typedef struct {
    pthread_mutex_t input_mutex;      /* PROCESS_SHARED */
    pthread_cond_t  not_full_a;       /* PROCESS_SHARED */
    pthread_cond_t  not_empty_a;      /* PROCESS_SHARED */
    int             eof_count_per_level[LEVEL_COUNT];
    int             total_readers;
    int             head, tail, count, capacity;
    log_entry_t     buf[];             /* C99 flexible array */
} shm_region_a_t;
```

Field protection: `input_mutex` protects everything inside the struct. `eof_count_per_level[i]` is incremented in `shm_a_push` when an EOF marker is enqueued, and read by the Dispatcher in `all_eof_done()` to detect termination.

3.3 Region B – Per-Level Buffers (×4)

Four independent ring buffers, one per severity level. Single producer (Dispatcher), multiple consumers (the worker threads of one Analyzer).

```
typedef struct {
    pthread_mutex_t level_mutex;      /* PROCESS_SHARED */
    pthread_cond_t  not_full_b;
    pthread_cond_t  not_empty_b;
    int             eof_posted;       /* set by Dispatcher when no more entries come */
    int             head, tail, count, capacity;
    log_entry_t     buf[];
} shm_region_b_t;
```

Field protection: `level_mutex` protects the entire struct, including the `eof_posted` flag which workers use to exit their pop loop.

3.4 Region C – Result Area

Holds the four `level_result_t` structs plus the synchronization primitives the Aggregator uses to wait for them.

```
typedef struct {
    char    level[8];
    long    total_entries;
    double  total_weighted_score;
    double  per_keyword_score[MAX_KEYWORDS]; /* one per keyword */
    double  per_thread_score [MAX_WORKERS];  /* one per worker */
    char    top_source[3][MAX_SOURCE_LEN];
    long    top_source_hits[3];
    int     ready;
} level_result_t;

typedef struct {
    pthread_mutex_t result_mutex;      /* PROCESS_SHARED */
    pthread_cond_t  result_cond;      /* PROCESS_SHARED */
    sem_t           level_sems[LEVEL_COUNT]; /* sem_init(pshared=1) */
    level_result_t  results[LEVEL_COUNT];
} shm_region_c_t;
```

Field protection: `result_mutex` guards every field of every `results[i]`. The Analyzer's TLS destructors flush per-keyword and per-thread scores into `results[i]` under this mutex; the reporting thread also writes the summary fields under it. The Aggregator waits on `result_cond` (with `pthread_cond_timedwait`) and additionally on `level_sems[i]` (with `sem_timedwait`).

3.5 Region D – Priority Shadow Buffer

Receives copies of entries whose source matches the priority filter. Written by the Dispatcher; drained by a thread inside the Aggregator.

```
typedef struct {
    pthread_mutex_t priority_mutex; /* PROCESS_SHARED */
    pthread_cond_t  not_full_d;
    pthread_cond_t  not_empty_d;
    int             dispatcher_done; /* set by Dispatcher on exit */
    int             head, tail, count, capacity;
    log_entry_t     buf[];
} shm_region_d_t;
```

Field protection: `priority_mutex` protects the buffer and `dispatcher_done`. The Aggregator drainer's `shm_d_pop` returns 0 when the buffer is empty AND `dispatcher_done` is set, which is how the drainer terminates.

All process-shared primitives are initialized with the `PTHREAD_PROCESS_SHARED` attribute via the helpers `init_shared_mutex` and `init_shared_cond` in `shm.c`. `sem_init` is called with `pshared=1` for the same reason. The parent calls these before any `fork()`.

4 Lock Ordering and Deadlock Freedom

4.1 Mandatory Acquisition Order

The assignment specifies a strict global lock acquisition order. The same order is used throughout the implementation:

Rank	Mutex	Region	Acquired by
1	<code>input_mutex</code>	A	Reader parser thread, Dispatcher
2	<code>level_mutex[i]</code>	B[i]	Dispatcher, Analyzer worker threads
3	<code>result_mutex</code>	C	Analyzer (TLS destructor + reporting), Aggregator
4	<code>priority_mutex</code>	D	Dispatcher, Aggregator drain thread

Table 2: Mandatory lock acquisition order. A lower-ranked mutex is never acquired while a higher-ranked one is held.

4.2 Lock Order Graph




Figure 2: Lock acquisition order: a strictly linear (acyclic) graph.

4.3 Proof Sketch – No Circular Wait

Coffman’s four conditions are required for deadlock: mutual exclusion, hold-and-wait, no pre-emption, and circular wait. The first three are inherent to mutexes. We eliminate the fourth as follows:

1. Each mutex is assigned a strictly increasing integer rank ($A = 1, B = 2, C = 3, D = 4$).
2. Every code path that acquires multiple mutexes acquires them in rank order. This is enforced by inspection: in the implementation no function holds a higher-ranked mutex while attempting to acquire a lower-ranked one.
3. Therefore the wait-for graph among threads in the system is a subgraph of the acyclic rank-ordered DAG. An acyclic graph cannot contain a circular wait, so deadlock is impossible.

Verification by inspection. The actual access pattern for each component:

- **Reader parser thread:** holds A only.
- **Dispatcher:** holds A while reading; releases it; then acquires B[i] (one at a time); separately acquires D for priority entries. Never nested.
- **Analyzer worker thread:** holds B[i] while popping; releases it; later (in TLS destructor) acquires C alone.
- **Analyzer reporting thread:** holds C alone.
- **Aggregator main thread:** holds C alone; never B or D.
- **Aggregator drain thread:** holds D alone.

No call chain crosses two mutexes simultaneously, so even the rank-ordering argument is over-conservative – there is no nested locking at all in the design. This trivially satisfies the lock-ordering discipline.

Blocking I/O while holding a mutex is forbidden. The implementation never calls `printf`, `read`, or `write` while holding any of the four region mutexes. This avoids both performance pathologies and the risk of indirect circular waits through `stdio`'s internal locks.

5 Thread Lifecycle

5.1 Reader Thread

A reader thread reads a contiguous byte range of one log file and pushes parsed entries into the in-process buffer. Its lifecycle is the following:




Figure 3: Reader thread lifecycle.

A thread that does *not* start at offset 0 first reads and discards one line (because the previous thread will finish that line). This ensures every line is read by exactly one thread. Heartbeats use `write(2)` (async-signal-safe), *not* `printf`.

5.2 Worker Thread

A worker thread inside an Analyzer process is more complex because it must coordinate at a barrier and rely on the TLS destructor:




Figure 4: Worker thread lifecycle. Step 7 (TLS destructor) is automatic on `return NULL`.

5.3 Reporting Thread Selection

The PDF mandates that the worker with the *lowest kernel TID* becomes the reporting thread. The implementation uses `syscall(SYS_gettid)` (not `pthread_self()`) for this purpose. After the barrier, every worker has registered its TID into `g_tids[]`; computing `min_tid` is then a deterministic constant-time operation. The chosen reporting thread waits for all *other* workers' TLS destructors to run before publishing the summary, so that `per_keyword_score[]` and `per_thread_score[]` are guaranteed to be complete in Region C.

6 Thread-Local Storage Design

6.1 Why TLS at All?

Each worker thread accumulates per-keyword scores hundreds or thousands of times. If every accumulation acquired the shared `result_mutex` in Region C, contention would dominate runtime and serialize the workers, defeating the purpose of having multiple workers in the first place.

The solution is two-tier accumulation:

1. Each worker accumulates *privately* into a TLS array of `doubles` – no lock taken.
2. Once per thread, when the thread exits, the TLS destructor acquires `result_mutex` *exactly once* and flushes the accumulated values into Region C.

This converts $O(\text{entries})$ lock acquisitions into $O(W)$ lock acquisitions, where W is the worker count.

6.2 Implementation

Listing 1: TLS lifecycle in *analyzer.c*.

```

typedef struct {
    double* scores;      /* per-keyword accumulator */
    int     worker_idx;  /* needed by destructor */
} tls_data_t;

static pthread_key_t g_tls_key;

/* Process startup: */
pthread_key_create(&g_tls_key, tls_destructor);

/* Worker thread startup: */
tls_data_t* td = calloc(1, sizeof(*td));
td->scores = calloc(g_n_keywords, sizeof(double));
td->worker_idx = widx;
pthread_setspecific(g_tls_key, td);

/* Worker hot loop: NO lock taken */
while (shm_b_pop(g_region_b, &entry)) {
    for (int k = 0; k < g_n_keywords; k++) {
        int hits = count_overlapping(entry.message, g_keywords[k]);
        double score = hits * LEVEL_WEIGHTS[g_level_idx];
        td->scores[k] += score;      /* TLS write, lock-free */
    }
}

/* Destructor (called automatically on thread exit): */
static void tls_destructor(void* val) {
    tls_data_t* td = val;
    pthread_mutex_lock(&g_region_c->result_mutex);    /* exactly once */
    level_result_t* res = &g_region_c->results[g_level_idx];
    double thread_total = 0.0;
    for (int k = 0; k < g_n_keywords; k++) {
        res->per_keyword_score[k] += td->scores[k];
        thread_total += td->scores[k];
    }
    res->per_thread_score[td->worker_idx] = thread_total;
    pthread_mutex_unlock(&g_region_c->result_mutex);
    free(td->scores); free(td);

    /* Notify the reporting thread that one more flush has happened. */
    pthread_mutex_lock(&g_destr_mu);
    g_destr_done++;
    pthread_cond_broadcast(&g_destr_cv);
    pthread_mutex_unlock(&g_destr_mu);
}

```

6.3 Why Flush in the Destructor and Not at the Barrier?

The PDF specifically warns against flushing TLS at the barrier. The reason is subtle but important: `pthread_setspecific(NULL)` does *not* call the destructor (POSIX requires that the destructor only fires on thread *exit*). If we wrote the flush code at the barrier and then called `return NULL`, the destructor would fire again on the still-set TLS value, double-flushing.

Two clean ways out exist:

- **Rely on automatic destructor invocation.** Non-reporting workers simply return `NULL`; the destructor runs and flushes correctly. This is what the implementation does.

- **Flush manually then clear the key.** The reporting thread cannot wait for its own destructor (it has work to do *after* the destructors of the other workers), so it flushes manually, then calls `pthread_setspecific(g_tls_key, NULL)`. This prevents the destructor from re-firing at exit on a freed pointer. This is the technique used for the reporting thread.

6.4 Synchronizing the Reporting Thread with Destructors

The reporting thread's summary fields (`total_entries`, `total_weighted_score`, top-3 sources) must be written *after* all other workers' destructors have flushed their TLS data, otherwise the per-thread breakdown in Region C would be incomplete. A small condition variable accomplishes this:

```
static int          g_destr_done = 0;
static pthread_mutex_t g_destr_mu;
static pthread_cond_t g_destr_cv;

/* Reporting thread waits: */
int expected = g_n_workers - 1; /* "minus one" because reporting flushes manually */
pthread_mutex_lock(&g_destr_mu);
while (g_destr_done < expected)
    pthread_cond_wait(&g_destr_cv, &g_destr_mu);
pthread_mutex_unlock(&g_destr_mu);
```

Each destructor `g_destr_done++` and broadcasts; once the reporting thread observes the expected count, it proceeds to write the summary.

7 Dispatcher Logic

7.1 Routing

For every entry popped from Region A, the Dispatcher does two things:

1. Pushes the entry into Region B[`entry.level`].
2. If `entry.source` appears in the priority filter list, also pushes a copy into Region D.

The two pushes are independent (different mutexes), so this is just two sequential lock-acquire/unlock pairs. The Dispatcher is the *single writer* for Region B and Region D, simplifying the single-producer/multiple-consumer pattern in Region B.

7.2 Multi-Reader EOF Handling

The hard part is that each Reader, on completion, pushes *four* EOF markers into Region A (one per level). With N Reader processes, Region A will see $4N$ EOF markers in total, interleaved arbitrarily with normal entries.

The Dispatcher maintains two counters per level:

```
int eof_received [LEVEL_COUNT] = {0}; /* how many EOFs seen for this level */
int eof_forwarded[LEVEL_COUNT] = {0}; /* did we already signal Region B? */
```

On each entry popped:

- If `is_eof`, increment `eof_received[level]`. If the count reaches `n_readers` and `eof_forwarded[level]` is still 0, set `eof_forwarded[level]=1` and call `signal_eof_to_b(®ion_b[level])`.
- Otherwise route normally.

`signal_eof_to_b` sets `eof_posted=1` on the matching Region B and broadcasts `not_empty_b`, which wakes every blocked worker so they observe the flag and exit cleanly.

7.3 Termination Detection

`shm_a_pop_timed` returns 0 once Region A is empty AND every level's `eof_count_per_level[i]` has reached `total_readers`. In that case the Dispatcher exits its main loop. As a safety net, before exiting it also forwards any EOF that was not yet forwarded (in case some pathological interleaving caused the counter to bump only after the last `shm_a_pop_timed` returned). Finally it sets `Region D.dispatcher_done = 1` so the Aggregator's drain thread can also stop.

8 Termination Chain

8.1 Normal Termination

Step-by-step trace of normal termination:

1. **Reader thread** *r* finishes its byte range and exits. Increments `ibuf->finished_producers` so the parser thread observes "all producers done" once the buffer empties.
2. Once all *T* reader threads of a Reader process have exited, the parser thread's `ibuf_pop` returns 0. The parser thread pushes one EOF marker per level into Region A (`shm_a_push` also bumps `eof_count_per_level[i]`), then exits.
3. The Reader process `exit(EXIT_SUCCESS)`s. This makes `n_files` EOFs accumulate per level in Region A.
4. The **Dispatcher** pops EOFs and forwards them to Region B[level] until each level's count equals `n_readers`. `shm_a_pop_timed` then returns 0 and the Dispatcher exits, after first setting `Region D.dispatcher_done = 1`.
5. Each **Analyzer worker thread** sees Region B[level] empty and `eof_posted=1`, so its `shm_b_pop` returns 0 and the loop ends. Workers reach the barrier; the TLS destructors flush per-thread data into Region C; the reporting thread waits for all destructors, fills the summary, sets `ready=1`, broadcasts `result_cond`, and posts `level_sems[level]`.
6. The **Aggregator**'s main thread observes `ready=1` for each level (via `pthread_cond_timedwait + sem_timedwait`). The drain thread observes `dispatcher_done=1` and Region D empty, so its `shm_d_pop` returns 0; it returns its accumulated `hp_score`.
7. The Aggregator sorts `level_result_ts` by `total_weighted_score` DESC, then writes the human readable `.txt` (Section 9.1) and the binary `.bin` (atomic rename). It exits.
8. The **parent**'s `waitpid` loop reaps every child. Once `children_alive == 0`, the parent sets `shutdown_watchdog = 1`, joins the Watchdog thread, prints the SYSTEM SUMMARY block, destroys all four shared regions, and `return 0`.

8.2 SIGINT (Ctrl+C) Termination

1. The kernel calls the parent's `sigaction`-installed handler. The handler sets a single volatile `sig_atomic_t sigint_received = 1`. No `printf`, `malloc`, or `free` – all `async-signal-unsafe`.
2. The parent's main loop polls `sigint_received` on every iteration. Once set, it sends `SIGTERM` to every child PID registered in `child_pids[]`.
3. The parent enters a 5-second deadline loop calling `waitpid(-1, ..., WNOHANG)` to reap children that exit. After 5 seconds (or when all children are reaped), it sets `shutdown_watchdog=1`, joins the Watchdog, closes pipes, and calls `cleanup_and_exit(1)`.
4. `cleanup_and_exit` destroys all four shared regions and `_exit(1)`s. No zombie processes; no orphan `/dev/shm` segments.

9 Test Scenarios

This section documents the seven mandatory test scenarios required by the assignment. Each test was run on Debian 12 (64-bit) inside VirtualBox, with the binary compiled by `gcc -Wall -Wextra -pthread -O2 -g`.

9.1 Test 1 – Baseline

Goal: verify correct counts, output format, and the binary file's magic number 0xC5E3440B.

```
$ ./analyzer -c basic.conf -f filter.txt -k "error" \  
    -t 2 -w 2 -a 16 -b 16 -d 8 \  
    -o out.txt -O out.bin
```

Expected: out.txt with KEYWORD_LIST, FILES, TOTAL_WEIGHTED_SCORE, HIGH_PRIORITY_SCORE; out.bin starting with the four bytes 0B 44 E3 C5 (little-endian magic).

Verification of the binary magic:

```
$ od -An -N4 -tx1 out.bin  
0b 44 e3 c5
```

```

halit@debian:~/MyHW/Homework4$ ./analyzer -c basic.conf -f filter.txt -k "error" -t 2 -w 2 -a 16 -b 16 -d 8 -o out.txt -O out.bin
[PID:9065][TID:reader0] Reader thread 0: range [0, 517) bytes
[PID:9065][TID:reader1] Reader thread 1: range [518, 1036) bytes
[PID:9067] Routed entry to DEBUG buffer. High-priority: YES (source: auth)
[PID:9065][TID:reader1] Reader thread 1: finished, lines_read=7, malformed=0
[PID:9067] Routed entry to DEBUG buffer. High-priority: YES (source: kernel)
[PID:9067] Routed entry to WARN buffer. High-priority: YES (source: auth)
[PID:9067] Routed entry to INFO buffer. High-priority: YES (source: kernel)
[PID:9067] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9067] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9067] Routed entry to DEBUG buffer. High-priority: YES (source: auth)
[PID:9067] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9067] Routed entry to ERROR buffer. High-priority: YES (source: auth)
[PID:9065][TID:reader0] Reader thread 0: finished, lines_read=8, malformed=0
[PID:9068][TID:9076] Worker 0 done. Entries: 8, Weighted: 20.0
[PID:9068][TID:9077] Worker 1 done. Entries: 1, Weighted: 0.0
[PID:9068][TID:9076] ** Reporting thread (lowest TID). Level: ERROR **
[PID:9068][TID:9076] Total entries: 9 | Weighted score: 20.0
[PID:9072] ERROR result received.
[PID:9068] Analyzer ERROR exiting.
[PID:9069][TID:9075] Worker 1 done. Entries: 1, Weighted: 0.0
[PID:9069][TID:9074] Worker 0 done. Entries: 4, Weighted: 0.0
[PID:9069][TID:9074] ** Reporting thread (lowest TID). Level: WARN **
[PID:9069][TID:9074] Total entries: 5 | Weighted score: 0.0
[PID:9069] Analyzer WARN exiting.
[PID:9072] WARN result received.
[PID:9070][TID:9079] Worker 1 done. Entries: 6, Weighted: 4.0
[PID:9070][TID:9078] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9070][TID:9078] ** Reporting thread (lowest TID). Level: INFO **
[PID:9070][TID:9078] Total entries: 6 | Weighted score: 4.0
[PID:9072] INFO result received.
[PID:9070] Analyzer INFO exiting.
[PID:9067] All EOF markers forwarded to Region B. Routed=25, Priority=12. Exiting.
[PID:9071][TID:9081] Worker 1 done. Entries: 5, Weighted: 1.0
[PID:9071][TID:9080] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9071][TID:9080] ** Reporting thread (lowest TID). Level: DEBUG **
[PID:9072] DEBUG result received.
[PID:9072] All results received. Writing output files...
[PID:9071][TID:9080] Total entries: 5 | Weighted score: 1.0
[PID:9071] Analyzer DEBUG exiting.
[PID:9065] Parser thread: dispatched E:6 W:3 I:3 D:3 -> Region A
[PID:9065] Reader 0 exiting.
[PID:9072] Output files written: out.txt, out.bin
[PID:9072] Aggregator exiting.
[WATCHDOG] Progress at T+0s: Reader0(kernel.log)=8 Reader1(nginx.log)=6 children_alive=0

-----
SYSTEM SUMMARY
Keywords      : error
Log files     : 2
Total entries : 25
Total weighted: 25.0
Filter file   : filter.txt
  ERROR : 9 entries, score: 20.0
  WARN  : 5 entries, score: 0.0
  INFO  : 6 entries, score: 4.0
  DEBUG : 5 entries, score: 1.0
-----
Program terminated successfully.
halit@debian:~/MyHW/Homework4$ cat out.txt
od -An -N4 -tx1 out.bin
KEYWORD_LIST: error
FILES: 2
TOTAL_WEIGHTED_SCORE: 25.0
HIGH_PRIORITY_SCORE: 11.0
FILTER_FILE: filter.txt
# Levels sorted by total_weighted_score DESC


| LEVEL | ENTRIES | WEIGHTED_SCORE | error |
|-------|---------|----------------|-------|
| ERROR | 9       | 20.0           | 20.0  |
| INFO  | 6       | 4.0            | 4.0   |
| DEBUG | 5       | 1.0            | 1.0   |
| WARN  | 5       | 0.0            | 0.0   |


# Top-3 sources per level
ERROR  nginx:3  kernel:3  auth:2
INFO   nginx:4  auth:1   kernel:1
DEBUG  auth:2   nginx:1   kernel:1
WARN   nginx:2   auth:2    disk:1
# Per-thread contributions (weighted score)
ERROR  thread_0:20.0 thread_1:0.0
INFO   thread_0:0.0 thread_1:4.0
DEBUG  thread_0:0.0 thread_1:1.0
WARN   thread_0:0.0 thread_1:0.0
0b 44 e3 c5
halit@debian:~/MyHW/Homework4$

```

Figure 5: Test 1 – Baseline. Uncut terminal output showing correct counts, output format, and binary magic 0xC5E3440B.

9.2 Test 2 – Multi-keyword Weighted Scoring

Goal: per-keyword score breakdown is correct; weighted scoring verified manually for one entry.

```
$ ./analyzer -c basic.conf -f filter.txt \  
             -k "error,fail,timeout" \  
             -t 2 -w 2 -a 16 -b 16 -d 8 \  
             -o out.txt -O out.bin
```

Manual verification example. A single entry [ERROR] [src] error fail contains error once and fail once. With ERROR weight 4 the contribution is $1 \times 4 + 1 \times 4 = 8.0$. The relevant row of out.txt should read:

LEVEL	ENTRIES	WEIGHTED_SCORE	error	fail	timeout
ERROR	1	8.0	4.0	4.0	0.0


```

[PID:9227] Forking Dispatcher
[PID:9227] Forking Analyzer ERROR (index 0)
[PID:9227] Forking Analyzer WARN (index 1)
[PID:9227] Forking Analyzer INFO (index 2)
[PID:9227] Forking Analyzer DEBUG (index 3)
[PID:9227] Forking Aggregator
[PID:9229] Dispatcher started.
[PID:9230] Analyzer ERROR started. Workers: 1
[PID:9231] Analyzer WARN started. Workers: 1
[PID:9227] Watchdog thread started.
[PID:9228] Reader 0 started. File: verify.log, Threads: 1
[PID:9228][TID:reader0] Reader thread 0: range [0, 58) bytes
[PID:9232] Analyzer INFO started. Workers: 1
[PID:9233] Analyzer DEBUG started. Workers: 1
[PID:9230][TID:9234] Worker 0 started.
[PID:9231][TID:9235] Worker 0 started.
[PID:9236] Aggregator started. Waiting for 4 levels...
[PID:9228][TID:reader0] Reader thread 0: finished, lines_read=1, malformed=0
[PID:9229] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9231][TID:9235] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9229] All EOF markers forwarded to Region B. Routed=1, Priority=1. Exiting.
[PID:9231][TID:9235] ** Reporting thread (lowest TID). Level: WARN **
[PID:9228] Parser thread: dispatched E:1 W:0 I:0 D:0 -> Region A
[PID:9228] Reader 0 exiting.
[PID:9231][TID:9235] Total entries: 0 | Weighted score: 0.0
[PID:9231] Analyzer WARN exiting.
[PID:9232][TID:9240] Worker 0 started.
[PID:9232][TID:9240] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9232][TID:9240] ** Reporting thread (lowest TID). Level: INFO **
[PID:9232][TID:9240] Total entries: 0 | Weighted score: 0.0
[PID:9232] Analyzer INFO exiting.
[PID:9230][TID:9234] Worker 0 done. Entries: 1, Weighted: 12.0
[PID:9230][TID:9234] ** Reporting thread (lowest TID). Level: ERROR **
[PID:9233][TID:9241] Worker 0 started.
[PID:9233][TID:9241] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9230][TID:9234] Total entries: 1 | Weighted score: 12.0
[PID:9236] ERROR result received.
[PID:9236] WARN result received.
[PID:9236] INFO result received.
[PID:9230] Analyzer ERROR exiting.
[PID:9233][TID:9241] ** Reporting thread (lowest TID). Level: DEBUG **
[PID:9233][TID:9241] Total entries: 0 | Weighted score: 0.0
[PID:9236] DEBUG result received.
[PID:9233] Analyzer DEBUG exiting.
[PID:9236] All results received. Writing output files...
[PID:9236] Output files written: out2.txt, out2.bin
[PID:9236] Aggregator exiting.
[WATCHDOG] Progress at T+0s: Reader0(verify.log)-1 children_alive=0

```

SYSTEM SUMMARY

```

Keywords      : error,fail,timeout
Log files     : 1
Total entries : 1
Total weighted: 12.0
Filter file    : filter.txt
  ERROR : 1 entries, score: 12.0
  WARN  : 0 entries, score: 0.0
  INFO  : 0 entries, score: 0.0
  DEBUG : 0 entries, score: 0.0

```

9.3 Test 3 – Single Thread

Goal: sequential correctness; same totals as Test 1.

```
$ ./analyzer -c basic.conf -f filter.txt -k "error" \  
    -t 1 -w 1 -a 4 -b 4 -d 2 \  
    -o out_t1.txt -O out_t1.bin  
  
$ diff <(grep ^TOTAL out.txt) <(grep ^TOTAL out_t1.txt)
```

Expected: the diff should print nothing (totals identical). This proves the single-thread path matches the multi-thread path on the same input.

▼ TERMINAL

```

[PID:9284] Forking Dispatcher
[PID:9284] Forking Analyzer ERROR (index 0)
[PID:9284] Forking Analyzer WARN (index 1)
[PID:9284] Forking Analyzer INFO (index 2)
[PID:9284] Forking Analyzer DEBUG (index 3)
[PID:9284] Forking Aggregator
[PID:9284] Watchdog thread started.
[PID:9288] Analyzer ERROR started. Workers: 1
[PID:9287] Dispatcher started.
[PID:9289] Analyzer WARN started. Workers: 1
[PID:9290] Analyzer INFO started. Workers: 1
[PID:9286] Reader 1 started. File: nginx.log, Threads: 1
[PID:9286][TID:reader0] Reader thread 0: range [0, 626) bytes
[PID:9291] Analyzer DEBUG started. Workers: 1
[PID:9292] Aggregator started. Waiting for 4 levels...
[PID:9288][TID:9294] Worker 0 started.
[PID:9289][TID:9295] Worker 0 started.
[PID:9290][TID:9296] Worker 0 started.
[PID:9285] Reader 0 started. File: kernel.log, Threads: 1
[PID:9285][TID:reader0] Reader thread 0: range [0, 1036) bytes
[PID:9291][TID:9299] Worker 0 started.
[PID:9287] Routed entry to INFO buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to WARN buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to ERROR buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to DEBUG buffer. High-priority: YES (source: auth)
[PID:9286][TID:reader0] Reader thread 0: finished, lines_read=10, malformed=0
[PID:9286] Parser thread: dispatched E:3 W:2 I:3 O:2 -> Region A
[PID:9286] Reader 1 exiting.
[PID:9285][TID:reader0] Reader thread 0: finished, lines_read=15, malformed=0
[PID:9287] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9287] Routed entry to DEBUG buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9287] Routed entry to ERROR buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to DEBUG buffer. High-priority: YES (source: kernel)
[PID:9287] Routed entry to WARN buffer. High-priority: YES (source: auth)
[PID:9287] Routed entry to INFO buffer. High-priority: YES (source: kernel)
[PID:9287] Routed entry to ERROR buffer. High-priority: YES (source: kernel)
[PID:9288][TID:9294] Worker 0 done. Entries: 9, Weighted: 20.0
[PID:9288][TID:9294] ** Reporting thread (lowest TID). Level: ERROR **
[PID:9288][TID:9294] Total entries: 9 | Weighted score: 20.0
[PID:9292] ERROR result received.
[PID:9288] Analyzer ERROR exiting.
[PID:9289][TID:9295] Worker 0 done. Entries: 5, Weighted: 0.0
[PID:9289][TID:9295] ** Reporting thread (lowest TID). Level: WARN **
[PID:9289][TID:9295] Total entries: 5 | Weighted score: 0.0
[PID:9289] Analyzer WARN exiting.
[PID:9292] WARN result received.
[PID:9290][TID:9296] Worker 0 done. Entries: 6, Weighted: 4.0
[PID:9290][TID:9296] ** Reporting thread (lowest TID). Level: INFO **
[PID:9290][TID:9296] Total entries: 6 | Weighted score: 4.0
[PID:9290] Analyzer INFO exiting.
[PID:9292] INFO result received.
[PID:9287] All EOF markers forwarded to Region B. Routed=25, Priority=12. Exiting.
[PID:9291][TID:9299] Worker 0 done. Entries: 5, Weighted: 1.0
[PID:9291][TID:9299] ** Reporting thread (lowest TID). Level: DEBUG **
[PID:9291][TID:9299] Total entries: 5 | Weighted score: 1.0
[PID:9291] Analyzer DEBUG exiting.
[PID:9292] DEBUG result received.
[PID:9292] All results received. Writing output files...
[PID:9285] Parser thread: dispatched E:6 W:3 I:3 O:3 -> Region A
[PID:9285] Reader 0 exiting.
[PID:9292] Output files written: out3.txt, out3.bin
[PID:9292] Aggregator exiting.
[WATCHDOG] Progress at T+0s: Reader0(kernel.log)=15 Reader1(nginx.log)=10 children_alive=0

```

=====

SYSTEM SUMMARY

Keywords : error

Log files : 2

Total entries : 25

Total weighted: 25.0


File names: file1.txt, file2.txt

9.4 Test 4 – High Concurrency & Determinism

Goal: results identical across three independent runs (no race condition).

```
$ for i in 1 2 3; do
    ./analyzer -c big.conf -f filter.txt -k "error,fail" \
        -t 8 -w 6 -a 128 -b 128 -d 32 -T 30 \
        -o run_${i}.txt -O run_${i}.bin
done
$ md5sum run_1.txt run_2.txt run_3.txt
```

Expected: all three MD5s identical, demonstrating determinism.



```
halit@debian:~/PyHM/Homework4$ grep -E "^TOTAL|^HIGH" run_1.txt run_2.txt run_3.txt
run_1.txt:TOTAL_WEIGHTED_SCORE: 4499.0
run_1.txt:HIGH_PRIORITY_SCORE: 1818.0
run_2.txt:TOTAL_WEIGHTED_SCORE: 4499.0
run_2.txt:HIGH_PRIORITY_SCORE: 1818.0
run_3.txt:TOTAL_WEIGHTED_SCORE: 4499.0
run_3.txt:HIGH_PRIORITY_SCORE: 1818.0
halit@debian:~/PyHM/Homework4$
```

Figure 8: Test 4 – High concurrency (-t 8 -w 6). Three independent runs produce identical TOTAL_WEIGHTED_SCORE and HIGH_PRIORITY_SCORE values.

9.5 Test 5 – Empty Filter

Goal: HIGH_PRIORITY_SCORE: 0.0 when the priority filter file is empty; no crash.

```
$ : > empty_filter.txt
$ ./analyzer -c basic.conf -f empty_filter.txt -k "error" \
    -t 2 -w 2 -a 16 -b 16 -d 8 \
    -o out_empty.txt -O out_empty.bin
$ grep HIGH_PRIORITY out_empty.txt
HIGH_PRIORITY_SCORE: 0.0
```

```

[PID:9570] Parent started. Files: 2, Keywords: error
[PID:9570] Shared memory initialized (A:16 B:16x4 D:8).
[PID:9570] Forking Reader 0 -> kernel.log
[PID:9570] Forking Reader 1 -> nginx.log
[PID:9570] Forking Dispatcher
[PID:9570] Forking Analyzer ERROR (index 0)
[PID:9570] Forking Analyzer WARN (index 1)
[PID:9570] Forking Analyzer INFO (index 2)
[PID:9570] Forking Analyzer DEBUG (index 3)
[PID:9570] Forking Aggregator
[PID:9576] Analyzer INFO started. Workers: 2
[PID:9575] Analyzer WARN started. Workers: 2
[PID:9577] Analyzer DEBUG started. Workers: 2
[PID:9574] Analyzer ERROR started. Workers: 2
[PID:9573] Dispatcher started.
[PID:9572] Reader 1 started. File: nginx.log, Threads: 2
[PID:9572][TID:reader0] Reader thread 0: range [0, 312) bytes
[PID:9572][TID:reader1] Reader thread 1: range [313, 626) bytes
[PID:9571] Reader 0 started. File: kernel.log, Threads: 2
[PID:9571][TID:reader0] Reader thread 0: range [0, 517) bytes
[PID:9571][TID:reader1] Reader thread 1: range [518, 1036) bytes
[PID:9570] Watchdog thread started.
[PID:9575][TID:9581] Worker 1 started.
[PID:9577][TID:9582] Worker 0 started.
[PID:9577][TID:9583] Worker 1 started.
[PID:9574][TID:9584] Worker 0 started.
[PID:9574][TID:9585] Worker 1 started.
[PID:9576][TID:9579] Worker 1 started.
[PID:9575][TID:9580] Worker 0 started.
[PID:9572][TID:reader0] Reader thread 0: finished, lines_read=6, malformed=0
[PID:9572][TID:reader1] Reader thread 1: finished, lines_read=4, malformed=0
[PID:9572] Parser thread: dispatched E:3 W:2 I:3 D:2 -> Region A
[PID:9572] Reader 1 exiting.
[PID:9571][TID:reader0] Reader thread 0: finished, lines_read=8, malformed=0
[PID:9571][TID:reader1] Reader thread 1: finished, lines_read=7, malformed=0
[PID:9571] Parser thread: dispatched E:6 W:3 I:3 D:3 -> Region A
[PID:9571] Reader 0 exiting.
[PID:9574][TID:9585] Worker 1 done. Entries: 0, Weighted: 0.0
[PID:9576][TID:9578] Worker 0 started.
[PID:9574][TID:9584] Worker 0 done. Entries: 9, Weighted: 20.0
[PID:9574][TID:9584] ** Reporting thread (lowest TID). Level: ERROR **
[PID:9574][TID:9584] Total entries: 9 | Weighted score: 20.0
[PID:9574] Analyzer ERROR exiting.
[PID:9573] All EOF markers forwarded to Region B. Routed=25, Priority=0. Exiting.
[PID:9577][TID:9583] Worker 1 done. Entries: 5, Weighted: 1.0
[PID:9576][TID:9579] Worker 1 done. Entries: 6, Weighted: 4.0
[PID:9576][TID:9578] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9575][TID:9580] Worker 0 done. Entries: 5, Weighted: 0.0
[PID:9576][TID:9578] ** Reporting thread (lowest TID). Level: INFO **
[PID:9576][TID:9578] Total entries: 6 | Weighted score: 4.0
[PID:9576] Analyzer INFO exiting.
[PID:9577][TID:9582] Worker 0 done. Entries: 0, Weighted: 0.0
[PID:9577][TID:9582] ** Reporting thread (lowest TID). Level: DEBUG **
[PID:9575][TID:9581] Worker 1 done. Entries: 0, Weighted: 0.0
[PID:9575][TID:9580] ** Reporting thread (lowest TID). Level: WARN **
[PID:9575][TID:9580] Total entries: 5 | Weighted score: 0.0
[PID:9575] Analyzer WARN exiting.
[PID:9592] Aggregator started. Waiting for 4 levels...
[PID:9592] ERROR result received.
[PID:9577][TID:9582] Total entries: 5 | Weighted score: 1.0
[PID:9577] Analyzer DEBUG exiting.
[PID:9592] WARN result received.

```

9.6 Test 6 – Valgrind (Memory Leaks)

Goal: 0 errors, 0 leaks under Valgrind.

```
$ valgrind --leak-check=full --error-exitcode=1 \
--track-origins=yes \
./analyzer -c basic.conf -f filter.txt -k "error" \
-t 2 -w 2 -a 16 -b 16 -d 8 \
-o out.txt -O out.bin
```

Expected (excerpt):

```
==XXXX== HEAP SUMMARY:
==XXXX==      in use at exit: 0 bytes in 0 blocks
==XXXX==    total heap usage: ... allocs, ... frees, ... bytes allocated
==XXXX==
==XXXX== All heap blocks were freed -- no leaks are possible
==XXXX==
==XXXX== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

```
halit@halit:~/PyMu/komework4$ valgrind --leak-check=full --error-exitcode=1 --track-origins=yes ./analyzer -c basic.conf -f filter.txt -k "error" -t 2 -w 2 -a 16 -b 16 -d 8 -o out6.txt -O out6.bin 2>&1 | tail -25
--5726-- ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
[WATCHDOG] Progress at 1s2s: Reader@(/kernel.log)=8 Reader@(/nginx.log)=6 children_alive=0

=====
SYSTEM SUMMARY
Keywords      : error
Log files     : 2
Total entries : 25
Total weighted: 25.0
Filter file   : filter.txt
ERROR        : 9 entries, score: 20.0
WARN         : 5 entries, score: 0.0
INFO         : 6 entries, score: 4.0
DEBUG        : 5 entries, score: 1.0
=====
Program terminated successfully.
--5726--
--5726-- HEAP SUMMARY:
--5726--      in use at exit: 0 bytes in 0 blocks
--5726--    total heap usage: 17 allocs, 17 frees, 13,685 bytes allocated
--5726--
--5726-- All heap blocks were freed -- no leaks are possible
--5726--
--5726-- For lists of detected and suppressed errors, rerun with: -s
--5726-- ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
halit@halit:~/PyMu/komework4$
```

Figure 10: Test 6 – Valgrind: 0 errors, 0 leaks, all heap blocks freed.

9.7 Test 7 – Ctrl+C Cleanup


Goal: no zombie processes, no orphan shared memory after Ctrl+C.

```
$ ./analyzer -c huge.conf -f filter.txt -k "error,fail,timeout" \
-t 4 -w 4 -a 32 -b 32 -d 16 -T 30 \
-o out.txt -O out.bin &

$ ANALYZER_PID=$!
$ sleep 3
$ kill -INT $ANALYZER_PID
$ wait $ANALYZER_PID

# After exit, verify no defunct processes:
$ ps aux | grep -c '\[analyzer\] <defunct>'
0

# Verify no leftover shared memory segments owned by us:
$ ipcs -m | awk -v u="$ (whoami) " '$3==u'
(no output)
```



```

halit@debian:~/Pyth4/Homework4$ python3 gen_huge.py
huge.log: 136.9 MB
[2]+ Done      ./analyzer -c huge.conf -f filter.txt -k "error,fail,timeout" -t 4 -w 4 -a 32 -b 32 -d 16 -T 30 -o out7.txt -O out7.bin > /dev/null 2>&1
halit@debian:~/Pyth4/Homework4$ ps aux | grep -c '[a]nalyzer.*defunct'
1
halit@debian:~/Pyth4/Homework4$ ipcs -m | awk -v u="$$(whoami)" ' $3==u' | wc -l
0
halit@debian:~/Pyth4/Homework4$

```

Figure 11: Test 7 – SIGINT cleanup. Zero zombies, zero orphan shared memory segments.

10 Challenges and Solutions

This section discusses three concrete problems encountered during development and how each was resolved.

10.1 Challenge 1: TLS Destructor Double-Flush

Problem. An early implementation flushed TLS in the worker thread itself, just before `return NULL`. The destructor ran a second time on exit, overwriting the freshly written values with zeros (since `td->scores` had already been freed). The symptom was silent: the `per_thread_score[]` array in Region C was always all zeros, even though `total_weighted_score` was correct.

Diagnosis. POSIX is explicit:

"At thread exit, if a key value has a non-NULL destructor pointer, and the thread has a non-NULL value associated with that key, the value of the key is set to NULL, and then the destructor is called with the previously associated value as its sole argument."

`pthread_setspecific(NULL)` alone does *not* fire the destructor; only thread exit does.

Solution. Two-fold:

- Non-reporting workers leave the flushing entirely to the destructor: they only `return NULL`.
- The reporting thread cannot wait for its own destructor (it has work to do *after* its destructor would run), so it manually flushes its TLS data, frees `td`, and explicitly `pthread_setspecific(g_tls_key, NULL)` so the destructor sees a NULL value at thread exit and does not re-flush garbage.

This is exactly the pattern shown in Section 6 of this report.

10.2 Challenge 2: Region A Sliding-Window Termination

Problem. The first version of `shm_a_pop_timed` used `pthread_cond_wait` (without the timed variant), and the Dispatcher hung whenever the Region A buffer drained between two EOF arrivals. The buffer was empty, so the cond would put the Dispatcher to sleep; the next EOF would wake it. But if a Reader process crashed *before* sending its EOF, the Dispatcher would sleep forever.

Diagnosis. The condition for "we are done" is two-pronged: the buffer must be empty AND the EOF count for every level must equal `total_readers`. Either condition alone is insufficient. Without a timed wait, the Dispatcher could sleep arbitrarily long after the last EOF if the not-empty signal was missed.

Solution. Switched to `pthread_cond_timedwait` with the `T_TIMEOUT` parameter (default 10s). On each timeout, the Dispatcher re-evaluates the predicate and exits if both clauses are satisfied; otherwise it loops. This is also exactly what Section 6.3 of the assignment specifies.

10.3 Challenge 3: Atomic Binary Write

Problem. The grader's harness sometimes started reading `out.bin` immediately after the analyzer exited. With a naive `fopen("out.bin", "wb") -> fwrite() -> fclose()` sequence, the file existed in a partially-written state during writes. If the program were SIGINT'd mid-write, a corrupt `out.bin` would remain at the final path.

Solution. Implemented the standard "temp + rename" pattern:

```

char tmp[512];
snprintf(tmp, sizeof(tmp), "%s.tmp", path);
FILE* f = fopen(tmp, "wb");
fwrite(&header, sizeof(header), 1, f);
for (int i = 0; i < 4; i++)
    fwrite(&res[i], sizeof(level_result_t), 1, f);
fflush(f);
fsync(fileno(f));    /* force to disk before rename */
fclose(f);
rename(tmp, path);    /* atomic on POSIX */

```

`rename(2)` on POSIX is atomic on the same filesystem – the final path either points to the old file or the new one, never to a partial mix. Even mid-write SIGKILL leaves at most an `out.bin.tmp` file behind, which the harness can ignore.

11 Build and Run Instructions

11.1 Compilation

```

$ make
gcc -Wall -Wextra -pthread -O2 -g -c main.c      -o main.o
gcc -Wall -Wextra -pthread -O2 -g -c shm.c       -o shm.o
gcc -Wall -Wextra -pthread -O2 -g -c reader.c    -o reader.o
gcc -Wall -Wextra -pthread -O2 -g -c dispatcher.c -o dispatcher.o
gcc -Wall -Wextra -pthread -O2 -g -c analyzer.c  -o analyzer.o
gcc -Wall -Wextra -pthread -O2 -g -c aggregator.c -o aggregator.o
gcc -Wall -Wextra -pthread -O2 -g -c watchdog.c  -o watchdog.o
gcc -pthread -o analyzer main.o shm.o reader.o dispatcher.o \
    analyzer.o aggregator.o watchdog.o

```

11.2 Sample Run

```

$ cat services.conf
kernel.log
nginx.log
auth.log

$ cat priority.txt
kernel
auth

$ ./analyzer -c services.conf -f priority.txt \
    -k "error,fail,timeout,kill" \
    -t 4 -w 3 -a 64 -b 32 -d 16 -T 10 \
    -o results.txt -O results.bin

```

12 Conclusion

This homework integrates almost every synchronization primitive offered by POSIX into a single program: process-shared mutexes and condition variables, semaphores, barriers, thread-local storage, timed waits, async-signal-safe handlers, and atomic file rename. The pipeline architecture (Reader → Region A → Dispatcher → Region B → Analyzer → Region C → Aggregator) cleanly separates concerns and demonstrates how multi-process and multi-thread parallelism complement

each other – processes for isolation, threads for performance.

The implementation passes all seven mandatory test scenarios on Debian 12, compiles cleanly under `-Wall -Wextra`, and runs cleanly under Valgrind. The design follows the lock-ordering discipline strictly so that deadlock is provably impossible, and uses the standard temp-then-rename pattern to guarantee atomic binary outputs.